

Collaboration Networks Process Book

Authors: Kevin Wall, Mike Liu, Will Usher

Timeline

Brainstorming

Networks




Figure 1: Dense subgraphs are replaced with small matrices

Early in the semester, Jean-Daniel Fekete gave a talk at the SCI institute on visualizing graphs using matrices in a method called [NodeTrix](#), introduced in a paper he coauthored. This gave us the idea to use hierarchical nodes to represent dense subgraphs instead of small matrices.

We looked at many different network/graph datasets available on the Stanford Network Analysis Project's website. However, these graphs were extremely large and were in general not very rich in information, only describing the graph itself and not telling much about the nodes.

After talking with the instructor, we were put on the trail of DBLP, which is a large database of research journals, articles, and authors. We decided we could make use this data to generate collaboration networks.

Proposal and Milestone Report

Our project proposal and milestone report are attached at the end of the PDF




Figure 2: Co-authorship network of scientists in field of network science



Figure 3: DBLP logo

Early Work

We already knew where our data was going to come from (the DBLP database), and we knew what we wanted to do with it, but a major roadblock to development was still getting the specific data we wanted to show (or at least a representative sample) in the format we wanted it in. As an example, we anticipated that scraping the author affiliations could be an expensive task, so we needed to know what websites to target as soon as possible. Without knowing what journals we were going to use, this task could not begin.

```
corr (CoRR); total url count: 90623
ieicet (IEICE Transactions); total url count: 17340
amc (Applied Mathematics and Computation); total url count: 14060
tit (IEEE Transactions on Information Theory); total url count: 13735
dm (Discrete Mathematics); total url count: 11854
tcs (Theor. Comput. Sci.); total url count: 10937
tsp (IEEE Transactions on Signal Processing); total url count: 10747
cacm (Commun. ACM); total url count: 10661
bioinformatics (Bioinformatics); total url count: 10370
eor (European Journal of Operational Research); total url count: 9743
eswa (Expert Syst. Appl.); total url count: 9001
iacr (IACR Cryptology ePrint Archive); total url count: 8712
tcom (IEEE Transactions on Communications); total url count: 8468
```

Figure 4: Output of journal analysis tool

In order to find these journals, we began making tools to analyze the data. The first such tool was a simple python script that collected endpoint URLs for a sample of publications within every journal in the DBLP database and printed out a their distribution, as well as which journals possessed the most links to publication databases. This told us where the articles were hosted and which journals had enough links to justify turning them into networks (we were concerned about getting author affiliations for all authors).




Figure 5: tog collaboration network

The next tool we made was a script that generated a GML file describing a collaboration network drawn from a inputted set of journals. This allowed for us to see the collaboration networks by opening the file with the graph visualization program Gephi. Now we could actually begin to judge potential datasets. Some networks were too dense, some weren't dense enough. Some had interesting structures we wished to visualize.

This script was also useful because it formed a basis for future scripts. It showed how we could read the DBLP database, demonstrated the need for a streaming XML parser due to the immense size of the database, and introduced an object-oriented model of Journals, Articles, and Authors. This made further analysis and processing much easier because we could simply access lists of class instances with useful per-instance information such as links to other instances. In addition, once we had a standard representation, it enabled better work parallelization.

After looking at various datasets and online publication libraries and applying our own biases towards certain journals, we settled on looking at journals that were published by ACM and hosted on the [digital library](#) and published by IEEE where we could use their API to get affiliation information.




Figure 6: Early visualization (1/3)

In order to visualize these datasets however, we needed a file format useful for communicating everything we new about the data, including information




Figure 7: Early visualization (2/3)




Figure 8: Early visualization (3/3)

we computed offline, to the Javascript that actually creates and controls the visualization.

This meant creating a new script that did essentially the same thing as the DBLP to gml script, but instead outputted JSON files. Now, finally, we could begin development of the visualization itself.




Figure 9: Test graph

Meanwhile, we were also developing a method of finding dense subgraphs within the networks we were generating. After doing some research, we found a approximate algorithm for finding dense subgraphs whose time complexity was linear [citation]. We implemented this, and were able to begin generating JSON files that described clusters in the data.

[discussion of early development of visualization]




Figure 10: SIGPLAN collaboration network with curated subgraph circled in red

Final Work

It became clear that many of our datasets were simply too large to visualize given our current methods (and were thus outside our scope). In order to reduce their size, we needed new tools to filter the data. This gave rise to two new developments. First, we modified the densest subgraph script to, after finding the densest subgraphs, find all the nodes connected to those subgraphs, and output the resulting graph as a modified version of the inputted JSON file. Second, we created a script that moved between our JSON file format and gml. This allowed us not only to visualize our graphs with Gephi, but also edit them with Gephi and then turn them back into JSON files. This allowed us to create curated collaboration networks with a manageable number of nodes.

Due to some positive early tests on the clustering implementation, some significant issues with it slipped through into late in development. What we discovered was that we were generating clusters of nodes that were unconnected and/or had a large number of loosely connected nodes. After analyzing the results and the implementation, we discovered that our algorithm was running into situation where the cluster we were looking for were narrowly losing to large collections of several clusters (in terms of density), especially when dealing with low densities. Our solution was to increase our minimum density requirements and to implement a special subgraph selection step that prefers smaller clusters if there exists close alternatives. You can see all this in figure 11; on the left you see the system attempting to bundle several disconnected clusters. On the right, the clusters now get their own bundles and more accurately reflect the structure of the graph.

[TODO MIKE: MORE discussion of final development of visualization]

In our final visualization we use a force directed graph with node bundles, the




Figure 11: Bad clustering of TIST collaboration network (left), good clustering (right)

nodes are the authors from the curated datasets and we bundle those within the clusters that were computed previously. This allows our Javascript code to remain relatively light weight as it doesn't need to perform a lot of heavy computation. Additionally we use the list of an author's or journal's articles from DBLP to show a listing on the side where the user can go read a journal or get more details about its authors.

Overview and Motivation

Collaboration Networks is an interactive visualization of author collaboration networks drawn from several ACM journals. It demonstrates a way to effectively visualize graphs that contain dense subgraphs without breaking with the visual language of nodes and edges.


Related Work

Our work is most influenced by [NodeTrix](#), which presents a method of managing dense subgraphs by using small matrices to represent them. Matrices provide a compact and informative way of communicating the connections in a dense subgraph, and by aggregating the connections flowing out of the subgraph, the resulting node-edge diagrams is much more visually manageable.

CS6630 Project: Collaboration Networks

members: Kevin Wall, Mike Liu, Will Usher

Select a Journal: SIGPLAN Notices



SIGPLAN Notices

Number of Articles:
1810

Years in Database:
1974 – 2013

Publications:

Challenges in Type Systems Research.

Authors:

• Martin Odersky

Year Published:

1997

<http://doi.acm.org/10.1145/251595.251607>

PADL '00: Workshop on Practical Aspects of Declarative Languages.

Authors:

• Enrico Pontelli
• Vitor Santos Costa

Year Published:

2000

Distribution of Articles Over Time




Figure 12: Screenshot of Collaboration Networks visualization

Questions

The main question we wanted to answer is whether dense subgraphs could be specially visualized to make the overall graph more readable while maintaining the language of nodes and edges, both macroscopically (in terms of the whole graph), and microscopically (in terms of the dense subgraph itself).

Data

Our main source of data was the DBLP database, which contains a large amount of data on various academic journals, articles, and authors. This data is in a structured format and available for download and so was easily accessible to us and did not require any special effort beyond filtering it down. Once it had been transformed into collaboration networks, however, the resulting networks required some more involved processing. This processing included finding dense subgraphs, connected subgraphs, and manual cleanup, removing connected subgraphs that were too small or too dense to be interesting or useful for visualization.

In addition to DBLP, we also scraped author affiliations from publication databases that were linked by the DBLP entries. This proved to be fairly involved as well as inconsistent in the case of some of the websites, and as a result, we ended up using only one of our initial prospects to avoid needing to support every different site layout. This limited the journals we could use to just those published by the ACM but was not so limiting as to be problematic.

Exploratory Data Analysis

Our main tool for viewing our data before we had a custom visualization method was Gephi. We used this software extensively to judge the qualities of the prospective collaboration network, as well as the quality of our densest subgraph implementation. There were many graphs that we viewed with Gephi and were able to immediately determine that we did not want to continue working with them either because they were far too dense, large, or not dense enough, their largest subgraphs only being a handful of nodes. We also used Gephi to evaluate experimental methods of simplifying dense graphs.

These methods in the end did not result in new visualizations due to the size of the graph still being an issue.

Of course, we also used our own visualizations to explore the data. This helped us determine our limits. Before our initial visualizations, we did not know either our performance limitations or visual space limitations. This information helped guide our data acquisition and methods themselves.




Figure 13: SIGPLAN collaboration network with computed clusters manually drawn on



Figure 14: TOG collaboration network unrestricted (left) and restricted to two author articles (right)



Figure 15: SIGPLAN collaboration network unrestricted (left) and restricted to two-to-three author articles (right)

Design Evolution

Before we had decided fully on networks, we created a initial concept of how we could visualize one. In figure 16, you can see our ideas beginning to take shape. Large-scale known grouping information is visualized as large dotted circles or tightly fitting shapes (We used both in order to compare the different methods). Clusters are grouped using solid black circles and the internal edges are not rendered (this doesn't show in most clusters because this visualization was only intended to get ideas across). When a node is selected, edges are shown (if hidden) and highlighted using red if the edges leave a cluster, and yellow if they are internal.

We discussed many different ideas on how to visualize the graph data that would help us discover interesting information amount the collaborations in the dataset. We thought about clustering data based on each node's affiliation, journal type, authors, or the year published to show different networks. We had a few different ideas on what method of visualizing the data would be most useful. They are the matrices method from the talk, voronoi diagram, and bubble set. We used Gephi to help us validate some ideas and settled on a hierarchy network graph could be an intuitive way to discover insights behind how researchers collaborated and published their findings.

We were considering using a bubble set library but found the library was pretty buggy and would require too much time to fix it. Later we found some D3.js examples with parts of components that we are looking at for our visualization like convex hulls (to show grouping of nodes), force-directed graphs, and bundling nodes (to show/hide dense clusters).

Convex hulls were useful for showing nodes in clusters, but not that good at displaying the network of the data. A force-directed graph was good for showing




Figure 16: First design concept

both clusters and the network but it don't help to simplify a hierarchical network graph data. Overall, we found that bundling nodes would be a good way to interact with our data and simplify dense clusters. It not only shows how authors are grouped, but also presents the network of authors based on their collaborations.

We learned that D3.js is not very performant when dealing with large amounts of data from, as a result we needed to reduce the size of data to be displayed, while keeping the interesting parts of the network. It was impossible to make data understandable. Therefore, We reduced the data by clustering authors within a journal and filtering nodes based on the density of clusters. Our visualization design went through a few iterations until we ended up on something we were happy with. We were thinking about adding different shape for distinguish clusters, but it made the visualization harder to understand. We only use the same color to show nodes are in the same cluster. We added summary of journals in our database when no journal is selected and an index overview to greet the user with initially. The summary gives an idea what data is available in our database to explore. In addition, we made the summary's bubble chart expand some when a moused over. Additionally the network of an author's collaboration with others is something we thought would be useful to explore, so this view was added as well. This view would help us to discover the possible reasons behind their collaborations, such as working at the same university or lab. We also have our side bar which displays all the available publications filter to be more interactive so the user can quickly select an author or doi of publications to view. We also have brush on a histogram that shows the amount of publications for each year for the selected journal and a user could selection the years they are intersted in finding and look at the network of authors by cutting the dataset down with the filter on the period of time.

Implementation

Our final visualization consists of three linked views. The main view is a graph showing the authors in the journal and collaborations between them where dense subgraphs of authors (those who frequently collaborate) are collapsed into bundles. A sidebar is shown which gives a more detailed summary of what's being shown in the main view, such as a list of papers in the journal or list of paper's by a specific author. Finally we provide a histogram of the papers published in the journal or by the author each year which can be brushed to select only papers within some time range. This brushing updates the main and side bar views to show just the papers and data within that selection.

In the main view we show a few different force directed graph layouts. The view greeting the user is a summary of the journals available in our database, where each journal is a node and they're sized by the number of papers in the journal. This will let them get a rough overview of which journals we have and what their contents might be like, e.g. old vs. young journals, those with many publications

etc. The user can then select a journal either by double clicking the node or selecting the journal from the dropdown menu.

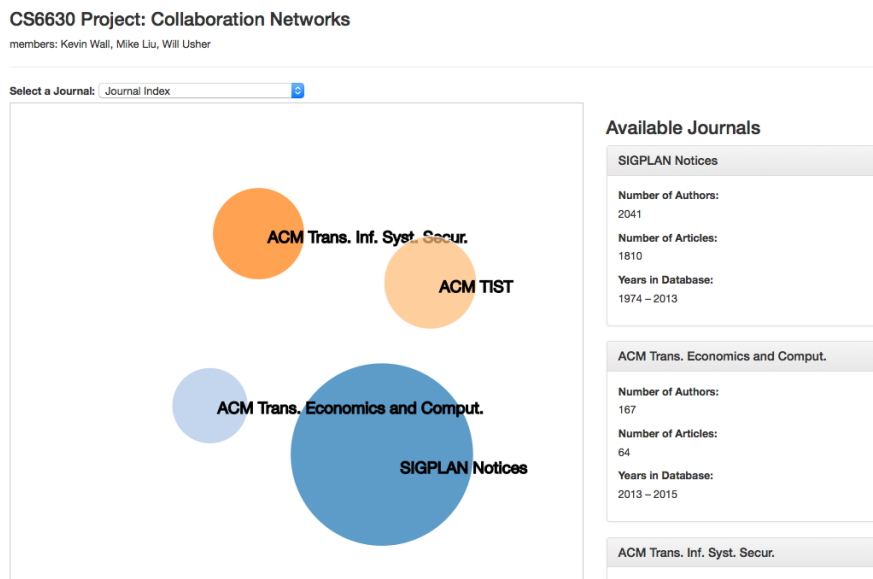


Figure 17: Landing view showing overview of journals

When a journal is selected we show a slightly curated network of authors in the journal and group those who frequently collaborate into clusters and bundle them as a single node. The curation is done to reduce the size of the dataset to more interesting clusters and subgraphs. In the journal view the clusters bundled as single larger nodes can be expanded by double clicking them to view the hidden dense collaboration network. The bundling of nodes helps keep the graph from becoming extremely cluttered and unreadable while also directing attention to interesting clusters (since the bundled nodes are bigger).


To keep the clusters visually grouped once expanded we color the nodes of the group the same color and draw a transparent convex hull around the nodes. Thus even for extremely dense and hard to read clusters the idea of “this is a dense collaborating group” is conveyed effectively. It’s easy to see that they’re all in the same convex hull and the mess of edges connecting all the authors to each other keeps the clump together. Additionally the edge between two clusters or authors is used to convey how frequently they collaborate by making it thicker in the case of more collaborations.

In the journal view the sidebar shows a list of the papers published in the journal along with some more details about the journal such as number of articles and

CS6630 Project: Collaboration Networks

members: Kevin Wall, Mike Liu, Will Usher

Select a Journal:



Year Published:
2000

A study of compiler techniques for multiple targets in compiler infrastructures.

Authors:

- Dai Gullian
- Zhang Suqing
- Tian Jinlan
- Jiang Weidu

Year Published:
2002

<http://doi.acm.org/10.1145/571727.571735>

Automatic construction of incremental LR(1)-parsers.

Authors:

- Dashing Yeh
- Uwe Kastens

Year Published:
1988

<http://doi.acm.org/10.1145/43895.43899>

Distribution of Articles Over Time




Figure 18: Full set of views showing a journal

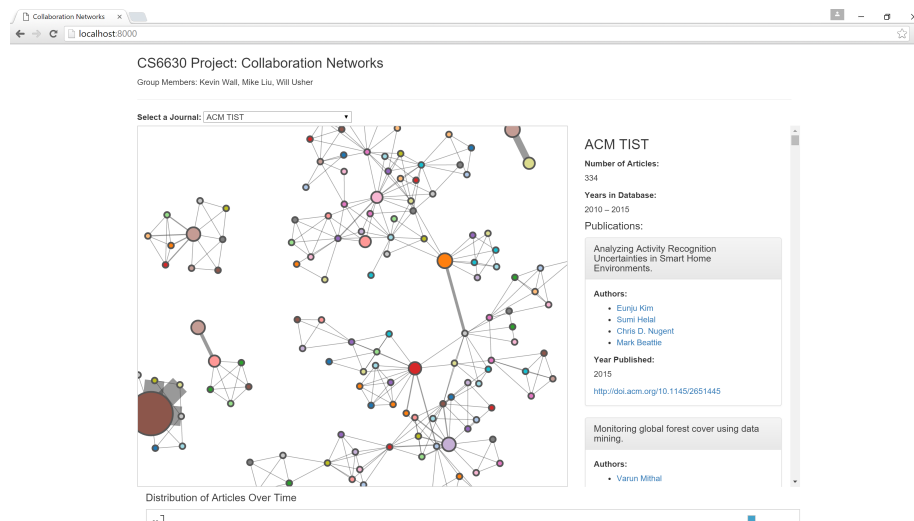


Figure 19: Unexpanded journal bundles

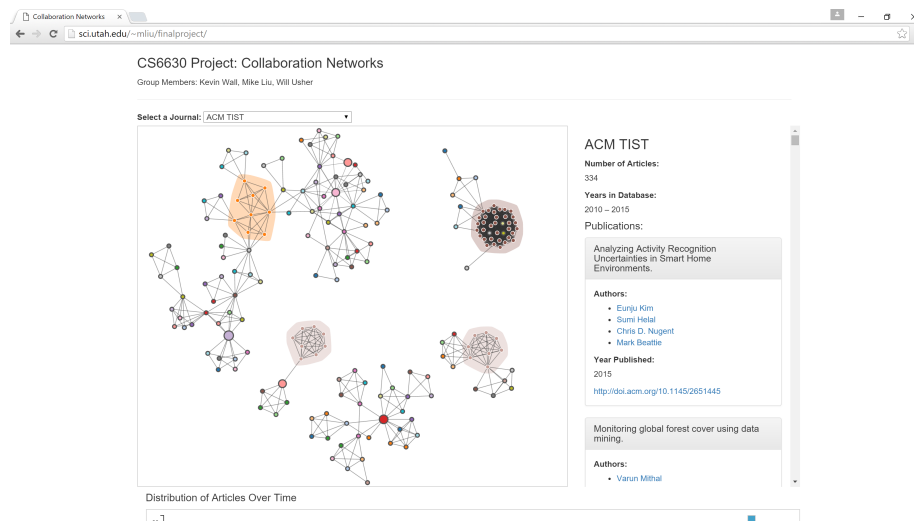


Figure 20: Expanded journal bundles

years of publications in the database. The article listing shows the title of the article, its authors, the year it was published and a DOI url to read the paper. The histogram view at the bottom shows the distribution of articles over time and can be brushed to select a subregion of time to view. This time filtering is applied both to the graph in the main view and to the list of articles in the sidebar so we can cut the view down to just collaborations over a few years instead of all time.

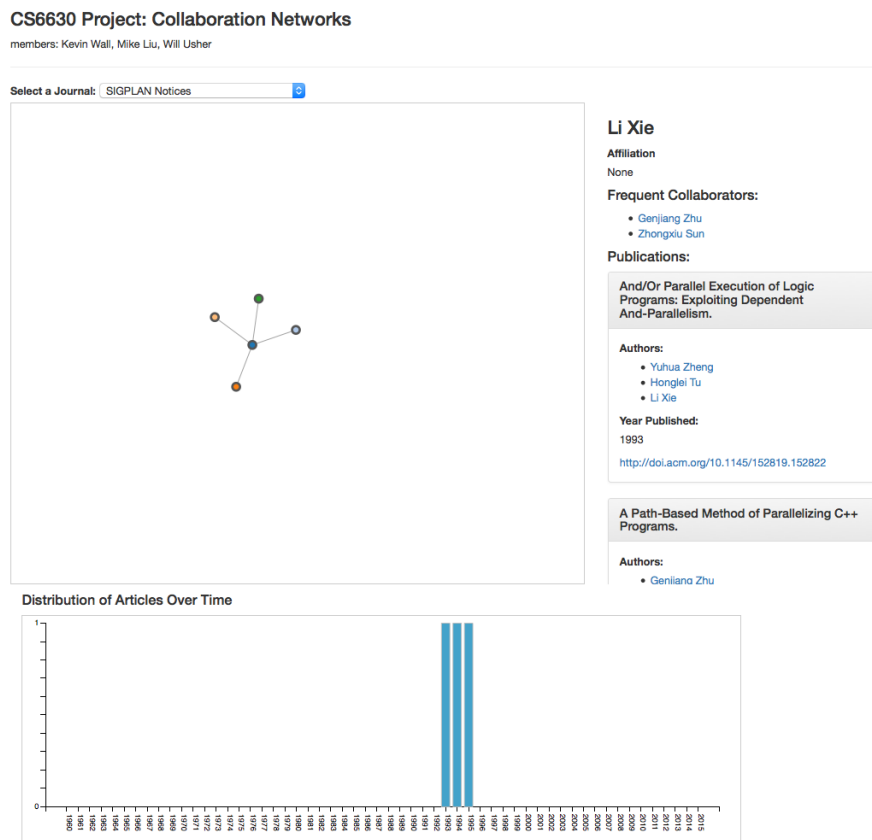


Figure 21: Author detail view

From this view the user can select an author to learn more about by either clicking the node in the graph or clicking their name on the list of publications. In the author view we show a graph of only those who have collaborated with them and in the side bar list their publications, affiliation and so on. Like in the journal view we also show the distribution of articles over time which the user can brush along to see who the author collaborated with over time.

Evaluation




Figure 22: SIGPLAN visualization with expanded bundle(left), with closed bundle(right)

One of the things we were worried about when coming up with this visualization was that dense clusters would always refer to articles with many authors. It turned out that this behavior was common in our final visualization (although we believe that this could still be solved with more advanced clustering methods). However this wasn't necessarily a bad thing which was surprising. Simplifying these dense clusters resulting from articles with many authors still gets rid of visual noise, and the edge aggregation also reveals interesting results. In Fig. 23 for instance, you can see that closed bundle visualization is cleaner while also communicating that there are a lot of collaborations between the blue and red bundles. There is a bug that we weren't able to fix in time is the display of edges would change back to default width's size when panning the view.

The brushing functionality could be better, especially if we recomputed clusters on the fly for the filtered dataset. This would allow for better exploration of subsets of the whole data which might be interesting to explore as well.

There are a couple additional features that would be nice to have such as highlight the node when mouse is over on an author in the side bar and highlight the year in the histogram and authors' node when mouse is over an article in the side bar. This would make the visualization even more connected.




Figure 23: TIST visualization with expanded bundle(left), with closed bundle(right)